

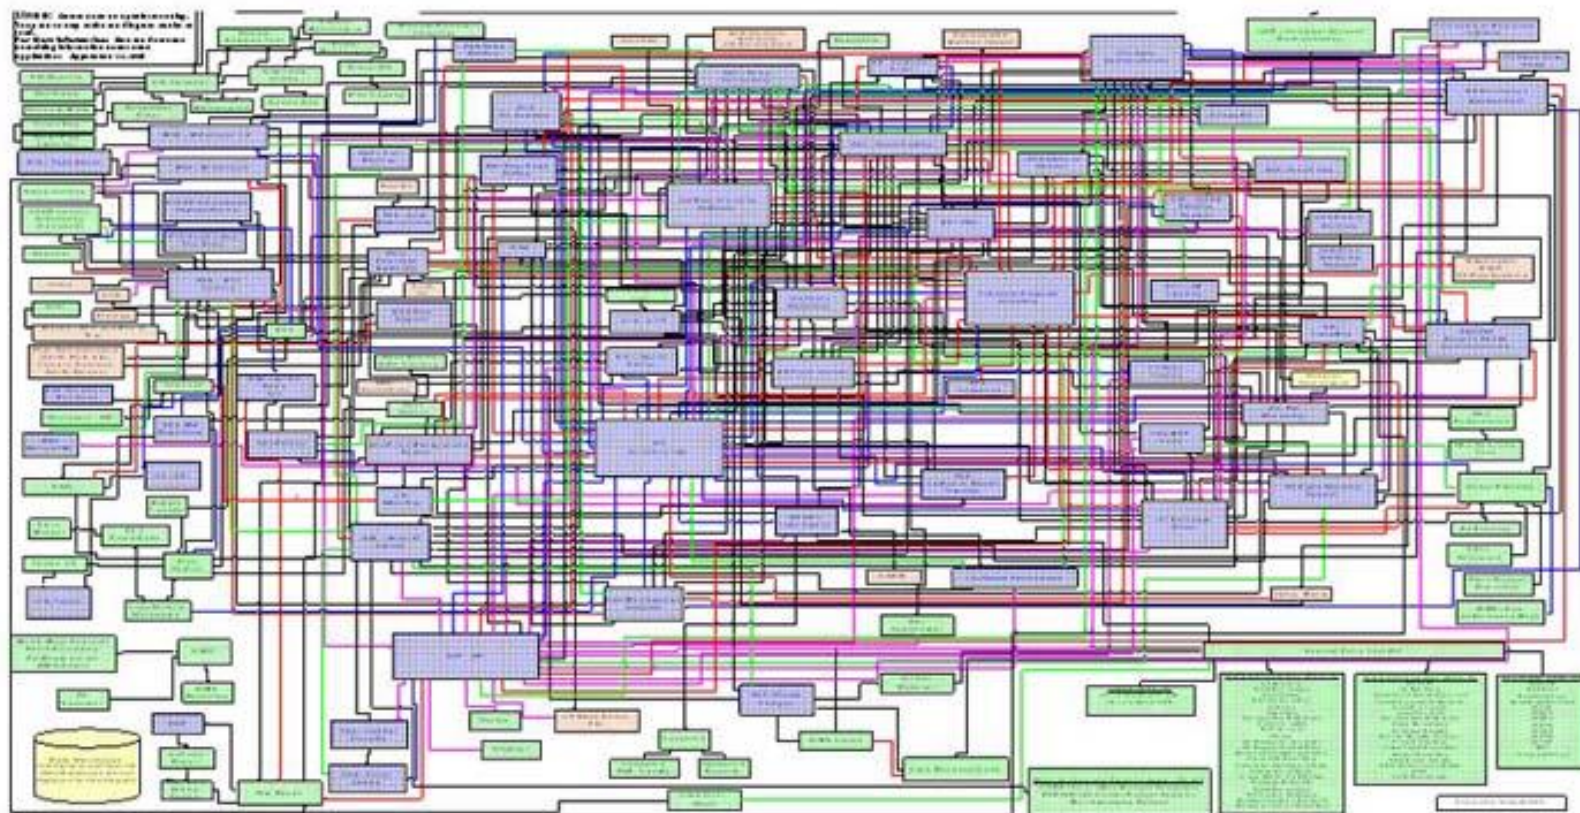


THE CASE FOR CHAOS TESTING

Peter Lamar
VP Cloud & Developer Relations
Peter.Lamar@Softwareag.com
peterlamar@gmail.com
[@ptlamar](https://twitter.com/ptlamar)

 **software** AG


COMPLEXITY



COMPLEXITY


Its easy to build software that becomes complex

- Fight complexity one line at a time



COMPLEXITY

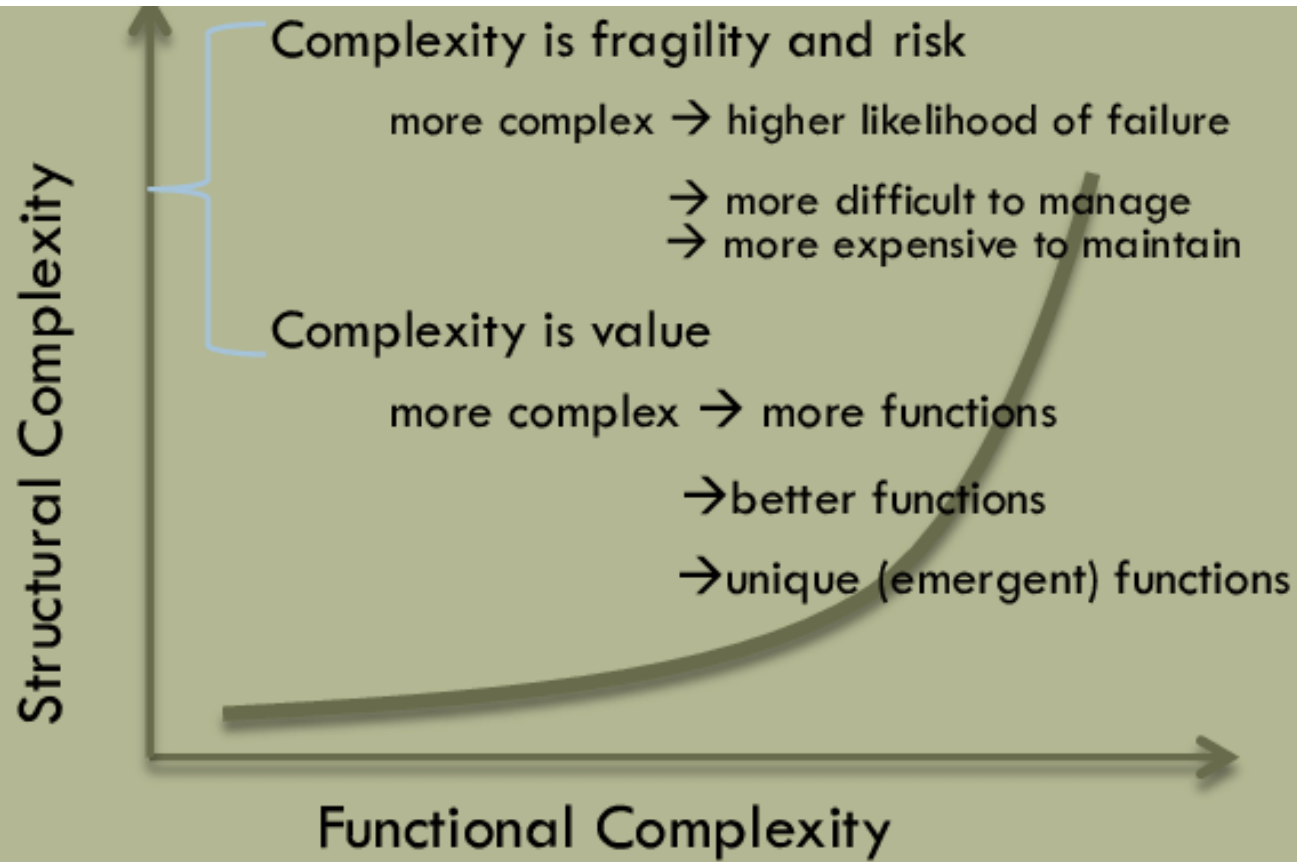
However, even the best systems grow in complexity over time



amazon.com



NETFLIX



Systems will grow in complexity, which will raise the likelihood of failure unless steps are taken to mitigate and manage complexity

COMPLEXITY

COMPLEXITY

Q. Do more professional, larger and better funded teams have less failure? i.e. Imposter syndrome?

A. Lets find out at
<https://outage.report>

(*Screenshot 8/22/19)

Recent Outages

Ongoing

Nitrado

Received 17 reports, originating from United States of America, Canada, Mexico, Germany, Switzerland and 2 more countries



Today

Microsoft Servers

Received 24 reports, originating from United States of America, Canada, Australia, Republic of Colombia, United Kingdom and 3 more countries



Hulu

Received 5 reports, originating from United States of America, Mexico, Norway



Minecraft

Received 23 reports, mostly originating from United States of America – Binghamton, Portland, Hartford, Providence, Kent and 10 more cities



Hunt: Showdown

Received 15 reports, originating from Germany, Russia, Brazil, Slovakia, Republic of France and 7 more countries



Gmail

Received 15 reports, originating from United States of America, Canada, Venezuela, Italy, Australia and 3 more countries



Verizon

Received reports, mostly originating from Dalton, Georgia, United States of America



Spotify

Received 11 reports, originating from United States of America, Republic of France, Canada, Brazil, Germany and 3 more countries



Nitrado

Received 4 reports, originating from Canada, United States of America, United Kingdom



Adobe Creative Cloud

Received 7 reports, originating from United States of America, Republic of Chile, Brazil, Germany, Republic of Austria and 1 more countries



COMPLEXITY

Q. Wow, all that failure must be expensive! Is there any way we can be more confident we are not going to fail?



A. Lets have controlled experiments where we try to identify weakness before it crashes the system.

COMPLEXITY

Sure, but 'controlled experiments' is a lame name. Lets call it 'Chaos Engineering'! Way cooler



COMPLEXITY



Chaos engineering is like a vaccine, which injects a small amount of virus to build immunity

CHAOS ENGINEERING

The harder it is to disrupt the steady state, the more confidence we have in the behavior of the system. If a weakness is uncovered, we now have a target for improvement before that behavior manifests in the system at large.

-<https://principlesofchaos.org>

CHAOS ENGINEERING — EXPECTED BENEFITS

Less downtime, better user experience



Less alarms and alerts (i.e. burnout) to
Operations/SRE/Development teams



More productivity from less unplanned outages



Spreading knowledge of application to the team

CHAOS ENGINEERING

1. Define the normal/steady state of the system (monitor system and business metrics)

Hypothesis that steady state will continue in both control and experiment groups



2. Pseudo-randomly inject faults (kill containers, network, etc) simulating real world events

Try to disprove hypothesis looking for difference in control and experiment groups

GAMEDAY EXAMPLE — PER SERVICE

Failure Scenario	Experiment	Scoping	Signals/Metrics	Abort Conditions
SAMPLE: Application server latency	Latency	100ms - 1000ms – 2500ms	Service Availability, On call paging	SLA breach, RPS threshold reach

Expected Outcome	Actual Outcome	Bugs
Application should still be available, but slower	Total app failure	Fallback did not occur

CNCF LANDSCAPE

Chaos Engineering



LIVE DEMO - SIMPLE CONTAINER EXPERIMENTS



PUMBA

Container Chaos experiment with Pumba

* Network delays!

LIVE DEMO — ADVANCED CONTAINER EXPERIMENTS

More advanced Chaos with Chaos Blade

* More Network Delays!



CHAOS ENGINEERING — SIMPLIFIED REVIEW

01

Have a hypothesis
and identify control
and experimental
group

02

Use real-world
events & limit scope

03

Make it as real as
possible, ideally
Production

04

Look for differences
in steady state
between control
and experimental
group

Test

Test failure of a stack component to identify resiliency of individual components



Test

Test failure for an Availability Zone/Region (Rack if on premise) to identify group component failover resiliency




Test

Test slowdown of key services to identify dependency slowdown/intermittent failure

EXAMPLE EXPERIMENTS

LIVE DEMO — CLUSTER EXPERIMENTS



Litmus

Test cluster failure with Litmus

ADDITIONAL TIPS

Start Small

Production if possible, or close to it

Minimize blast radius

Have an emergency stop

QUESTIONS?

Peter Lamar

VP Cloud & Developer Relations

Peter.Lamar@Softwareag.com

peterlamar@gmail.com

@ptlamar

